

Grouping and Sorting with cuDF

In this notebook you will be introduced to grouping and sorting with cuDF, with performance comparisons to Pandas, before integrating what you learned in a short data analysis exercise.

Objectives

By the time you complete this notebook you will be able to:

- Perform GPU-accelerated group and sort operations with cuDF

Imports

```
In [1]: import cudf
import pandas as pd
```

Read Data

We once again read the UK population data, returning to timed comparisons with Pandas.

```
In [2]: %time gdf = cudf.read_csv('./data/pop_1-04.csv', dtype=['float32', 'str', 'str', 'float32'])
CPU times: user 2.08 s, sys: 600 ms, total: 2.68 s
Wall time: 2.68 s
```

```
In [3]: %time df = pd.read_csv('./data/pop_1-04.csv')
CPU times: user 26 s, sys: 3.69 s, total: 29.7 s
Wall time: 29.7 s
```

```
In [4]: gdf.dtypes
```

```
Out[4]: age      float32
sex         object
county      object
lat         float32
long        float32
name        object
dtype: object
```

```
In [5]: gdf.shape
```

```
Out[5]: (58479894, 6)
```

```
In [6]: gdf.head()
```

```
Out[6]:
```

	age	sex	county	lat	long	name
0	0.0	m	Darlington	54.533638	-1.524400	Francis
1	0.0	m	Darlington	54.426254	-1.465314	Edward
2	0.0	m	Darlington	54.555199	-1.496417	Teddy
3	0.0	m	Darlington	54.547905	-1.572341	Angus
4	0.0	m	Darlington	54.477638	-1.605995	Charlie

Grouping and Sorting

Record Grouping

Record grouping with cuDF works the same way as in Pandas.

cuDF

```
In [7]: %%time
counties = gdf[['county', 'age']].groupby(['county'])
avg_ages = counties.mean()
print(avg_ages[:5])
```

```

              age
county
Warrington    40.888416
Reading       35.868777
Derbyshire    42.913279
East Sussex   44.757385
Northumberland 44.626919
CPU times: user 64 ms, sys: 12 ms, total: 76 ms
Wall time: 75.2 ms
```

Pandas

```
In [8]: %%time
counties_pd = df[['county', 'age']].groupby(['county'])
avg_ages_pd = counties_pd.mean()
print(avg_ages_pd[:5])
```

```

              age
county
Barking And Dagenham    33.056845
Barnet                  37.629770
Barnsley                 41.201061
Bath And North East Somerset 39.822837
Bedford                 39.715300
CPU times: user 3.9 s, sys: 994 ms, total: 4.89 s
Wall time: 4.87 s
```

Sorting

Sorting is also very similar to Pandas, though cuDF does not support in-place sorting.

cuDF

```
In [9]: %time gdf_names = gdf['name'].sort_values()
print(gdf_names[:5]) # yes, "A" is an infrequent but correct given name in the UK, acc
print(gdf_names[-5:])
```

CPU times: user 1.65 s, sys: 7.59 ms, total: 1.66 s

Wall time: 1.66 s

26850 A

154537 A

165578 A

211428 A

236972 A

Name: name, dtype: object

58060377 Zyrh

58289490 Zyrh

58363665 Zyrh

58388727 Zyrh

58394184 Zyrh

Name: name, dtype: object

Pandas

This operation takes a while with Pandas. Feel free to start the next exercise while you wait.

```
In [10]: %time df_names = df['name'].sort_values()
print(df_names[:5])
print(df_names[-5:])
```

CPU times: user 1min 43s, sys: 1.62 s, total: 1min 45s

Wall time: 1min 45s

10811041 A

17931460 A

5060367 A

1842288 A

24866365 A

Name: name, dtype: object

47008072 Zyrh

47953653 Zyrh

31838209 Zyrh

53669567 Zyrh

54557840 Zyrh

Name: name, dtype: object

Exercise: Youngest Names

For this exercise you will need to use both `groupby` and `sort_values`.

We would like to know which names are associated with the lowest average age and how many people have those names. Using the `mean` and `count` methods on the data grouped by `name`, identify the three names with the lowest mean age and their counts.

```
In [12]: name_groups = gdf[['name', 'age']].groupby('name')

name_ages = name_groups['age'].mean()
name_counts = name_groups['age'].count()

ages_counts = cudf.DataFrame()
ages_counts['mean_age'] = name_ages
ages_counts['count'] = name_counts

ages_counts = ages_counts.sort_values('mean_age')
ages_counts.iloc[:3]
```

```
Out[12]:
```

	mean_age	count
name		
Leart	34.911197	259
Luke-Junior	35.313725	255
Nameer	35.479675	246

Solution

```
In [13]: # %Load solutions/youngest_names
name_groups = gdf[['name', 'age']].groupby('name')

name_ages = name_groups['age'].mean()
name_counts = name_groups['age'].count()

ages_counts = cudf.DataFrame()
ages_counts['mean_age'] = name_ages
ages_counts['count'] = name_counts

ages_counts = ages_counts.sort_values('mean_age')
ages_counts.iloc[:3]
```

```
Out[13]:
```

	mean_age	count
name		
Leart	34.911197	259
Luke-Junior	35.313725	255
Nameer	35.479675	246

Please Restart the Kernel

```
In [ ]: import IPython
app = IPython.Application.instance()
app.kernel.do_shutdown(True)
```

Next

As part of our larger data science goal for this workshop, we will be working with data reflecting the entire road network of Great Britain. In the next notebook you will be exposed to additional cuDF techniques that you will use to transform columnar data into graph edge data that we will be using to construct a GPU-accelerated graph using the `cuGraph` library.